

# The Price of Inference: A Longitudinal Economic Analysis of Hierarchical LLM Credential Leakage

Tanishq S

*Dept. of Computer Science*

*Rajalakshmi Institute of Technology*

Chennai, India

stanishq07@gmail.com

Vignesh B

*Dept. of Computer Science*

*Rajalakshmi Institute of Technology*

Chennai, India

bvignesh1464@gmail.com

**Abstract**—The integration of Large Language Models (LLMs) into the software supply chain has fundamentally altered the nature of credential leakage. Unlike static secrets (e.g., database passwords), LLM API keys are *liquid economic assets*—direct bearers of computational reasoning and token quotas. This paper introduces the Theory of Economic Credential Stratification, identifying a critical divergence in how these assets are managed compared to their utility. Using CHRONOS, a custom-built forensic instrument, we analyzed exposed artifacts across GitHub. We report a paradoxical “Protection Gap”: Tier 1 (GPT-4) credentials—despite having high abuse potential for *LLMjacking* [5]—exhibit a mean survival time ( $t_{surv}$ ) of 48 hours in non-production artifacts, significantly longer than lower-value keys. Furthermore, we identify “Dataset Poisoning” as a systemic blind spot: valid credentials embedded in .jsonl and .parquet training files often persist indefinitely, becoming part of the model’s latent knowledge.

**Index Terms**—LLM Security, Credential Leakage, Supply Chain Security, Forensics, Economic Stratification

## I. INTRODUCTION

In the “Post-API Economy,” authentication tokens are currency. While Meli et al. (2019) characterized secret leakage in open-source ecosystems [1], the security community has largely treated secrets as binary entities: either valid or invalid. This view is obsolete. We posit that leakage dynamics are driven by **Credential Economics**. A leaked AWS key allows infrastructure abuse; a leaked OpenAI key allows *inference theft*—unwittingly subsidizing a malicious actor’s compute costs. **Contributions:**

- 1) **Economic Stratification:** A taxonomy classifying secrets by *economic clearance* (Tier 1 vs. Tier 2).
- 2) **The CHRONOS Framework:** An active, hierarchical probing methodology (Fig. 1).
- 3) **Discovery of the “Dataset Blindsight”:** Empirical evidence of “dark matter” leaks in large-scale ML data formats.

## II. RELATED WORK

Existing literature focuses on static analysis of secrets [1], [2] or adversarial extraction from model weights [4], [10]. However, a gap exists at the intersection of *Supply Chain Security* and *LLM Economics*. **Static Secret Scanning:** Tools like Trufflehog or GitGuardian excel at pattern matching but lack semantic understanding of the *value* of the asset. They

treat a GPT-4 key identical to a test database password. **Model Inversion:** Carlini et al. demonstrated that training data can be extracted from models. Our work in Section II-B inverts this threat: instead of extracting data from models, we find that models are often trained on data that contains valid credentials, creating a feedback loop of compromise.

## III. METHODOLOGY: THE CHRONOS FRAMEWORK

To quantify these dynamics, we architected **CHRONOS**, a longitudinal monitoring system designed for active economic probing.

### A. System Architecture

The system operates as a closed-loop active scanner. The *Discovery Module* identifies candidates, which are passed to the *Forensics Engine* for context extraction. The *HMCs* then probes the key’s value, while the *Lifecycle Daemon* tracks its survival.

### B. Hierarchical Model Clearance System (HMCs)

We reject the boolean validation model ( $V : K \rightarrow \{0, 1\}$ ). Instead, we implement a hierarchical probe Algorithm 1. **Definition 1 (Credential Asset Value).** Let  $K$  be a set of credentials. The asset value  $V(k)$  for  $k \in K$  is defined as the latent vector potential of the associated model  $M_k$  constrained by its quota  $Q_k$ :

$$V(k) = \mathcal{F}(M_k) \times Q_k$$

where Tier 1 represents  $M_k \in \{GPT-4, Claude-3-Opus\}$ .

## IV. THREAT MODELING: THE RACE TO REVOCATION

We model the lifecycle of a leak as a race condition between three agents: The *Leaker* (Developer), the *Attacker*, and the *Defender*.

### A. The Race Condition

The window of vulnerability,  $\Delta t$ , is defined as the time difference between the *Public Event* ( $t_{pub}$ ) and the *First Revocation Event* ( $t_{rev}$ ).

$$\Delta t = t_{rev} - t_{pub}$$

If the Attacker’s reaction time  $t_{atk} < \Delta t$ , the asset is compromised.




Fig. 1. The CHRONOS System Architecture.

### Algorithm 1: HMCS Probing Logic

---

```

Candidate Key  $k$  Tier  $T \in \{0, 1, 2, \emptyset\}$   $R_1 \leftarrow \text{POST} /v1/chat (\text{model}=\text{'gpt-4'})$ ;
if  $R_1.\text{status} == 200$  then
| return Tier 1 (Strategic);
end
if  $R_1.\text{status} == 429$  then
| return Tier 0 (Exhausted);
end
 $R_2 \leftarrow \text{POST} /v1/chat (\text{model}=\text{'gpt-3.5'})$ ;
if  $R_2.\text{status} == 200$  then
| return Tier 2 (Consumer);
end
return Invalid;
  
```

---




Fig. 2. The HMCS Decision Logic Flowchart.




Fig. 3. The Leak Race Condition (Detailed View).

## V. EMPIRICAL EVALUATION

*Analysis based on  $N=50,000$  artifacts.*

### A. The “Inverse Value” Paradox

Contrary to rational expectations, **Tier 1 (GPT-4) keys survive longer than Tier 2 keys** in `test/` directories.

- **Hypothesis:** High-value keys are often owned by senior developers or automated systems (CI/CD) that bypass standard pre-commit hooks.

- **Impact:** Attackers have longer dwell time on the most valuable assets.

#### B. The “Dataset Poisoning” Vector

We identified a significant cluster of valid keys within `training_data.jsonl` files. Because these files are large binary/text blobs, they are rarely scanned.




Fig. 4. Distribution of leak sources.

#### C. Longitudinal Survival Data

Our survival analysis reveals distinct decay rates for different credential tiers.




Fig. 5. Survival Decay Rates.

#### D. Data Rigor & Schema

To ensure reproducibility, all data was captured using a rigorous schema designed for forensic auditability.




Fig. 6. Entity Relationship Diagram (ERD).

## VI. DISCUSSION & REMEDIATION

### A. From Detection to Remediation

The era of *Passive Detection* is ending. Future systems must move to *Active Remediation*. We propose **Autonomous Security Agents (ASAs)** that, upon detecting a leak, instantaneously rotate the credential via the provider’s API.

### B. Policy Implications

Current compliance standards (SOC2, ISO 27001) are insufficient for LLM assets. We recommend:

- **Dynamic Quotas:** API keys should have floating usage limits ( $Q_{max}$ ) that adapt to their deployment context (dev vs. prod).
- **Dataset Sanitization:** Mandatory pre-training scans for credential patterns in ‘.jsonl’ corpora.

### C. Limitations

Our study was limited to GitHub public events. Private repositories and GitLab/Bitbucket ecosystems may exhibit different decay dynamics. Furthermore, we did not test “Canary Tokens,” focusing only on live functional credentials.

## VII. CONCLUSION

Project CHRONOS demonstrates that the *economics* of the leaked secret determine its lifecycle. By ignoring the “stratification” of access, we leave the door open to the most damaging class of AI-driven attacks.

## REFERENCES

- [1] Meli, M., et al. (2019). “How Bad Can It Git?” *NDSS*.
- [2] GitGuardian. (2024). “State of Secrets Sprawl 2024.”
- [3] Wiz Research. (2025). “The AI Security Gap.”
- [4] Carlini, N., et al. (2023). “Extracting Training Data from Large Language Models.” *USENIX Security*.
- [5] Lasso Security. (2024). “The Guide to LLMjacking.”
- [6] ReversingLabs. (2024). “The State of Software Supply Chain Security 2024.”
- [7] Kroll. (2024). “Secrets Management Report: The State of Cyber Defense.”
- [8] OWASP. (2025). “Top 10 for LLM Applications: System Prompt Leakage & Vector Weaknesses.”
- [9] Hoplon Infosec. (2024). “LLM Security: The Prompt Injection Threat.”
- [10] Carlini, N., et al. (2021). “Extracting Training Data from Large Language Models.”
- [11] SISA Infosec. (2024). “Secrets in the Crawl: Analysis of Common Crawl Data.”
- [12] IAPP. (2024). “Privacy Risks in LLM Training Data.”
- [13] Protecto.ai. (2024). “Automated PII Removal in GenAI Pipelines.”
- [14] Cloud Security Alliance. (2024). “Optimizing Secrets Management.”
- [15] IBM Security. (2024). “Cost of a Data Breach Report 2024.”
- [16] CrowdStrike. (2024). “2024 Global Threat Report.”
- [17] SentinelOne. (2024). “The Future of AI Security.”
- [18] Palo Alto Networks. (2024). “Unit 42 Cloud Threat Report, Vol 8.”
- [19] Forrester. (2024). “The State of Application Security, 2024.”

- [20] Prompt Security. (2024). “The Rise of Prompt Injection Attacks.”
- [21] Akeyless. (2024). “State of Secrets Management Survey Report.”
- [22] Gartner. (2024). “Security and Risk Management Summit Insights.”
- [23] ACM. (2024). “Proceedings of the 2024 ACM SIGSAC Conference.”
- [24] IEEE. (2024). “IEEE Symposium on Security and Privacy (SP).”
- [25] Usenix. (2024). “33rd USENIX Security Symposium.”